

```

import pandas as pd
from sklearn.cluster import KMeans
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler

# Step 1: Load the CSV files
customers = pd.read_csv('Customers.csv')
transactions = pd.read_csv('Transactions.csv')

# Step 2: Merge data
# Merge transactions with customer information
merged_data = pd.merge(transactions, customers, on='CustomerID',
                        how='inner')

merged_data = merged_data.drop(columns='ProductID')
merged_data

```

	TransactionID	CustomerID	TransactionDate	Quantity
TotalValue \				
0	T00001	C0199	2024-08-25 12:38:23	1
300.68				
1	T00112	C0146	2024-05-27 22:23:54	1
300.68				
2	T00166	C0127	2024-04-25 07:38:55	1
300.68				
3	T00272	C0087	2024-03-26 22:55:37	2
601.36				
4	T00363	C0070	2024-03-21 15:10:10	3
902.04				
..	...	...	...	...
.				
995	T00496	C0118	2024-10-24 08:30:27	1
459.86				
996	T00759	C0059	2024-06-04 02:15:24	3
1379.58				
997	T00922	C0018	2024-04-05 13:05:32	4
1839.44				
998	T00959	C0115	2024-09-29 10:16:02	2
919.72				
999	T00992	C0024	2024-04-21 10:52:24	1
459.86				
	Price	CustomerName	Region	SignupDate
0	300.68	Andrea Jenkins	Europe	2022-12-03
1	300.68	Brittany Harvey	Asia	2024-09-04
2	300.68	Kathryn Stevens	Europe	2024-04-04
3	300.68	Travis Campbell	South America	2024-04-11
4	300.68	Timothy Perez	Europe	2022-03-15

```

..      ...      ...      ...      ...
995  459.86      Jacob Holt  South America  2022-01-22
996  459.86  Mrs. Kimberly Wright  North America  2024-04-07
997  459.86      Tyler Haynes  North America  2024-09-21
998  459.86      Joshua Hamilton      Asia  2024-11-11
999  459.86      Michele Cooley  North America  2024-02-05

```

```
[1000 rows x 9 columns]
```

```
# Step 3: Feature Engineering
```

```
# Calculate total spent by each customer
```

```

customer_total_spent = merged_data.groupby('CustomerID')
['TotalValue'].sum().reset_index()
customer_total_spent.columns = ['CustomerID', 'TotalSpent']

```

```
# Calculate the average cart value by each customer
```

```

customer_avg_cart_value = merged_data.groupby('CustomerID')
['TotalValue'].mean().reset_index()
customer_avg_cart_value.columns = ['CustomerID', 'AvgCartValue']

```

```
# Calculate the number of transactions by each customer (purchase frequency)
```

```

customer_frequency = merged_data.groupby('CustomerID')
['TransactionID'].nunique().reset_index()
customer_frequency.columns = ['CustomerID', 'PurchaseFrequency']

```

```

customer_features = pd.merge(customer_total_spent,
customer_avg_cart_value, on='CustomerID')
customer_features = pd.merge(customer_features, customer_frequency,
on='CustomerID')

```

```
customer_features
```

```

      CustomerID  TotalSpent  AvgCartValue  PurchaseFrequency
0          C0001    3354.52    670.904000                5
1          C0002    1862.74    465.685000                4
2          C0003    2725.38    681.345000                4
3          C0004    5354.88    669.360000                8
4          C0005    2034.24    678.080000                3
..      ...      ...      ...      ...
194        C0196    4982.88    1245.720000                4
195        C0197    1928.65    642.883333                3
196        C0198     931.83    465.915000                2
197        C0199    1979.28    494.820000                4
198        C0200    4758.60    951.720000                5

```

```
[199 rows x 4 columns]
```

```

scaler = StandardScaler()
scaled_features =
scaler.fit_transform(customer_features[['TotalSpent', 'AvgCartValue',

```

```
'PurchaseFrequency']])
```

```
# Step 2: Elbow Method to find the optimal number of clusters
```

```
wcss = [] # List to store WCSS for each K
```

```
for k in range(2, 11): # Testing from 2 to 10 clusters
```

```
    kmeans = KMeans(n_clusters=k, random_state=42)
```

```
    kmeans.fit(scaled_features)
```

```
    wcss.append(kmeans.inertia_) # WCSS is stored in inertia_
```

```
# Plotting Elbow Curve
```

```
plt.figure(figsize=(8, 6))
```

```
plt.plot(range(2, 11), wcss, marker='o')
```

```
plt.title('Elbow Method For Optimal K')
```

```
plt.xlabel('Number of Clusters')
```

```
plt.ylabel('WCSS (Within-Cluster Sum of Squares)')
```

```
plt.show()
```

```
C:\Users\revat\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1429: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
You can avoid it by setting the environment variable
OMP_NUM_THREADS=1.
```

```
    warnings.warn(
```

```
C:\Users\revat\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1429: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
You can avoid it by setting the environment variable
OMP_NUM_THREADS=1.
```

```
    warnings.warn(
```

```
C:\Users\revat\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1429: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
You can avoid it by setting the environment variable
OMP_NUM_THREADS=1.
```

```
    warnings.warn(
```

```
C:\Users\revat\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1429: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
You can avoid it by setting the environment variable
OMP_NUM_THREADS=1.
```


```
    warnings.warn(
```

```
C:\Users\revat\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1429: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
You can avoid it by setting the environment variable
OMP_NUM_THREADS=1.
```

```
    warnings.warn(
```

```
C:\Users\revat\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1429: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
```

```
You can avoid it by setting the environment variable
OMP_NUM_THREADS=1.
warnings.warn(
C:\Users\revat\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1429: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
You can avoid it by setting the environment variable
OMP_NUM_THREADS=1.
warnings.warn(
C:\Users\revat\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1429: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
You can avoid it by setting the environment variable
OMP_NUM_THREADS=1.
warnings.warn(
C:\Users\revat\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1429: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
You can avoid it by setting the environment variable
OMP_NUM_THREADS=1.
warnings.warn(
```



```
from sklearn.metrics import davies_bouldin_score

optimal_k = 5

# Step 4: Perform K-means clustering with the chosen number of clusters
kmeans = KMeans(n_clusters=optimal_k, random_state=42)
customer_features['Segment'] = kmeans.fit_predict(scaled_features)

# Step 5: Calculate Davies-Bouldin Index and other metrics
db_index = davies_bouldin_score(scaled_features,
customer_features['Segment'])
print(f'Davies-Bouldin Index: {db_index}')

# Additional metrics (e.g., silhouette score, calinski harabasz score)
can also be calculated
from sklearn.metrics import silhouette_score, calinski_harabasz_score

sil_score = silhouette_score(scaled_features,
customer_features['Segment'])
```

```

calinski_score = calinski_harabasz_score(scaled_features,
customer_features['Segment'])

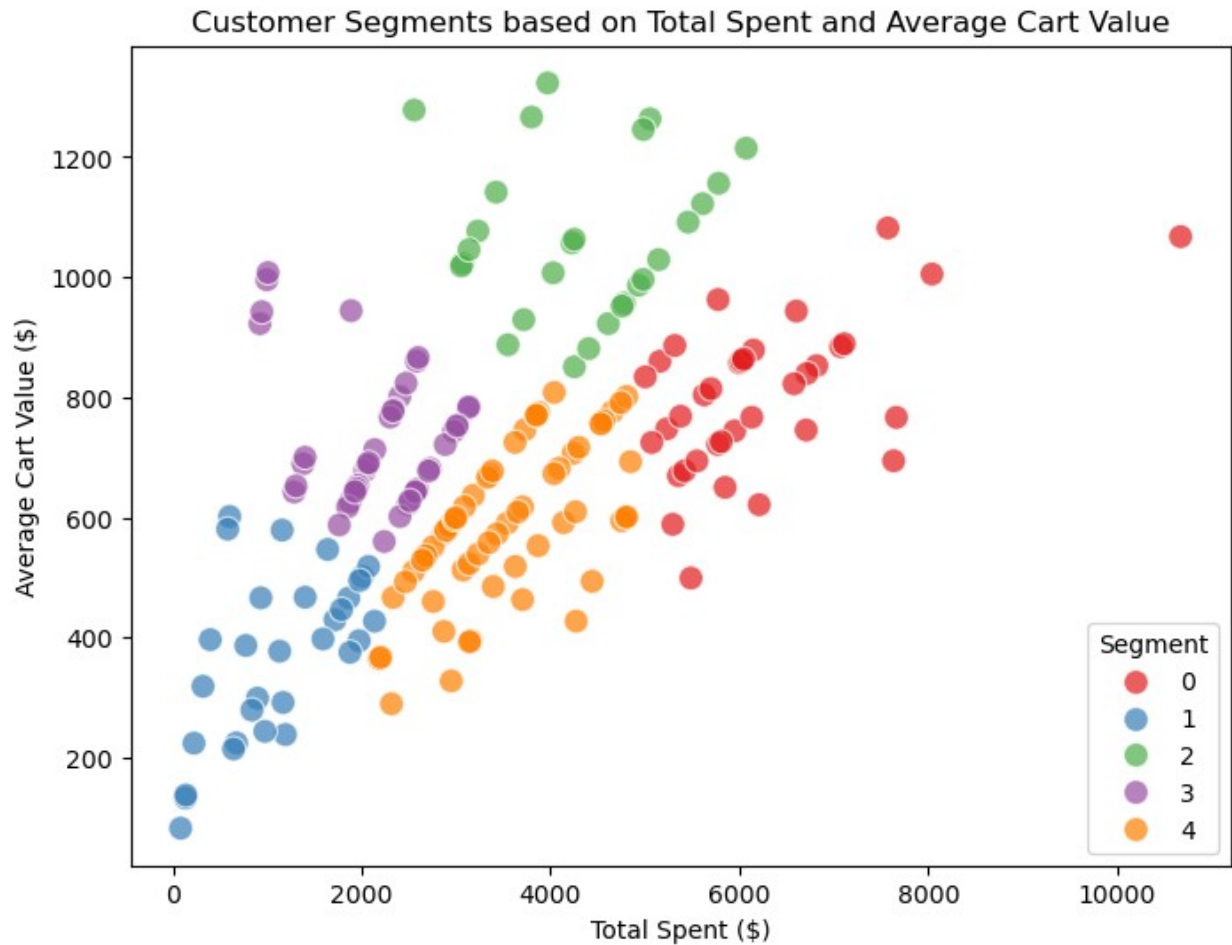
print(f'Silhouette Score: {sil_score}')
print(f'Calinski-Harabasz Score: {calinski_score}')

# Step 6: Visualizing the Clusters (2D scatter plot for better
visualization)
plt.figure(figsize=(8, 6))
sns.scatterplot(x=customer_features['TotalSpent'],
y=customer_features['AvgCartValue'], hue=customer_features['Segment'],
palette='Set1', s=100, alpha=0.7)
plt.title('Customer Segments based on Total Spent and Average Cart
Value')
plt.xlabel('Total Spent ($)')
plt.ylabel('Average Cart Value ($)')
plt.legend(title='Segment')
plt.show()

C:\Users\revat\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1429: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
You can avoid it by setting the environment variable
OMP_NUM_THREADS=1.
  warnings.warn(

Davies-Bouldin Index: 0.8524813520458036
Silhouette Score: 0.3535127066812944
Calinski-Harabasz Score: 143.70464826443802

```



```
# Step 2: Perform K-Means Clustering (K=4 as chosen from previous
code)
from sklearn.cluster import KMeans, AgglomerativeClustering
from sklearn.mixture import GaussianMixture

# Step 3: Perform Agglomerative Hierarchical Clustering (using ward
linkage)
agg_clust = AgglomerativeClustering(n_clusters=5, linkage='ward')
customer_features['Agglomerative_Segment'] =
agg_clust.fit_predict(scaled_features)

agg_db = davies_bouldin_score(scaled_features,
customer_features['Agglomerative_Segment'])

print(f"Agglomerative Davies-Bouldin Index: {agg_db}")

# Additional metrics (Silhouette Score and Calinski-Harabasz Score)
print("\nAdditional Metrics:")
print(f"Agglomerative Silhouette Score:
{silhouette_score(scaled_features,
customer_features['Agglomerative_Segment'])}")
```

```
print(f"Agglomerative Calinski-Harabasz Score:
{calinski_harabasz_score(scaled_features,
customer_features['Agglomerative_Segment'])}")

# Agglomerative Clusters
plt.figure(figsize=(8, 6))
sns.scatterplot(x=customer_features['TotalSpent'],
y=customer_features['AvgCartValue'],
hue=customer_features['Agglomerative_Segment'], palette='Set2', s=100,
alpha=0.7)
plt.title('Agglomerative Clusters (Total Spent vs Average Cart
Value)')
plt.xlabel('Total Spent ($)')
plt.ylabel('Average Cart Value ($)')
plt.legend(title='Segment')
plt.show()
```

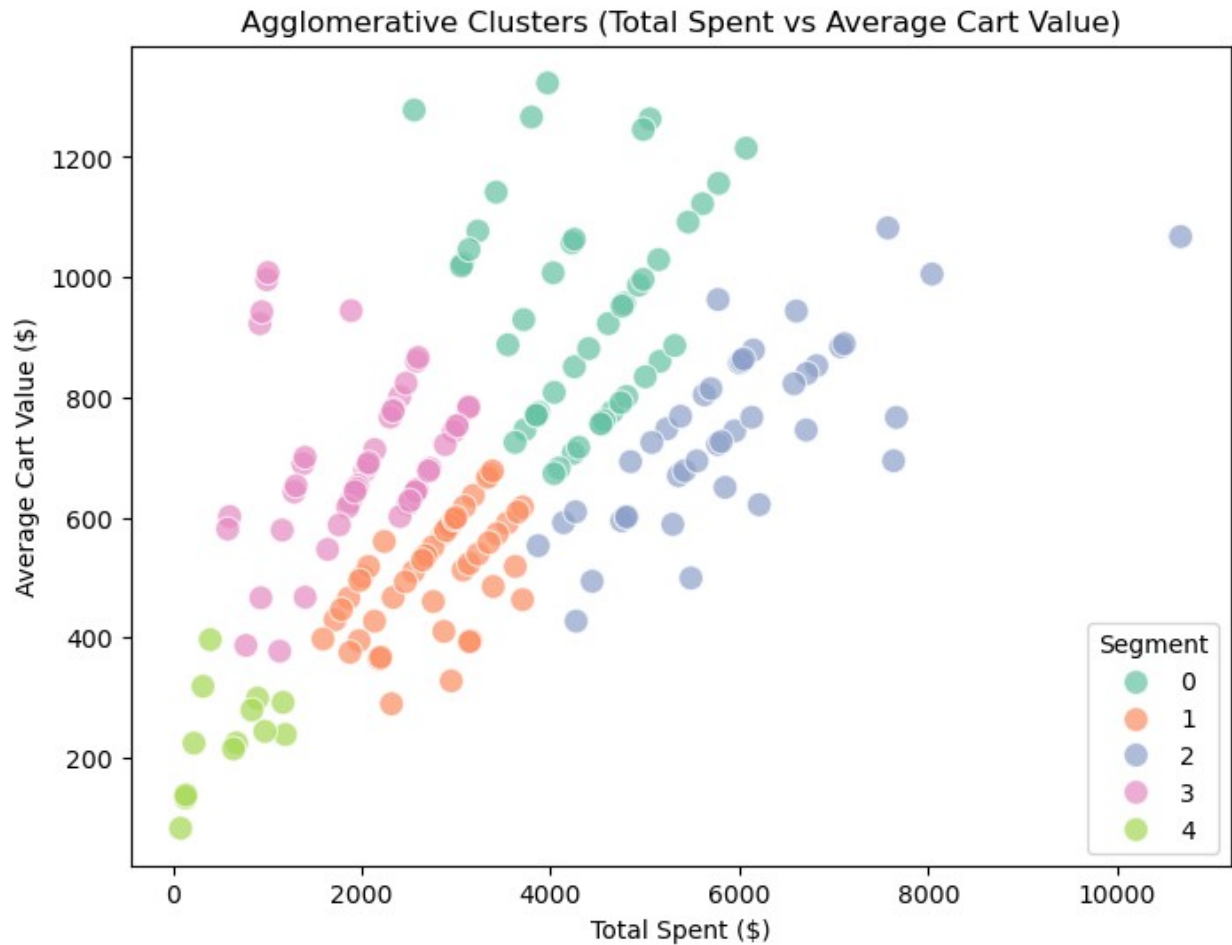
Agglomerative Davies-Bouldin Index: 0.8851081906022241

Additional Metrics:

Agglomerative Silhouette Score: 0.34265137202494406

Agglomerative Calinski-Harabasz Score: 130.61439024040507





```
# K-Means Clustering
import os
os.environ["OMP_NUM_THREADS"] = "1"

from sklearn.cluster import KMeans

kmeans = KMeans(n_clusters=3, random_state=42)
customer_features['KMeans_Segment'] = kmeans.fit_predict(
    customer_features[['TotalSpent', 'AvgCartValue',
    'PurchaseFrequency']]
)

# Agglomerative Clustering
from sklearn.cluster import AgglomerativeClustering

agglomerative = AgglomerativeClustering(n_clusters=3)
customer_features['Agglomerative_Segment'] =
agglomerative.fit_predict(
    customer_features[['TotalSpent', 'AvgCartValue',
    'PurchaseFrequency']]
)
```

```
)  
print(customer_features.shape)
```

```
C:\Users\revat\anaconda3\Lib\site-packages\sklearn\cluster\  
_kmeans.py:1429: UserWarning: KMeans is known to have a memory leak on  
Windows with MKL, when there are less chunks than available threads.  
You can avoid it by setting the environment variable  
OMP_NUM_THREADS=1.  
  warnings.warn(
```

```
(199, 7)
```